

Simple Linear Regression

```
In [2]: #Name: Yash Pravin Gadbail  
#Roll no. : 35  
#Sec: 3rd A  
#Sub : ET 1  
#Date:05/10/2024
```

```
In [4]: # Aim: To Perform Simple Linear Regression
```

```
In [6]: import pandas as pd
```

```
In [106]: df=pd.read_csv("C:\\Users\\OneDrive\\Desktop\\Salary.csv")
```

```
In [ ]:
```

```
In [108]: df.head()
```

```
Out[108]:
```

| | YearsExperience | Salary |
|---|-----------------|--------|
| 0 | 1.1 | 39343 |
| 1 | 1.3 | 46205 |
| 2 | 1.5 | 37731 |
| 3 | 2.0 | 43525 |
| 4 | 2.2 | 39891 |
| 5 | 2.9 | 56642 |
| 6 | 3.0 | 60150 |
| 7 | 3.2 | 54445 |
| 8 | 3.2 | 64445 |
| 9 | 3.7 | 57189 |

```
In [110]: df.tail()
```

```
Out[110]:
```

| | YearsExperience | Salary |
|----|-----------------|--------|
| 30 | 11.2 | 127345 |
| 31 | 11.5 | 126756 |
| 32 | 12.3 | 128765 |
| 33 | 12.9 | 135675 |
| 34 | 13.5 | 139465 |

In [112]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 35 entries, 0 to 34
Data columns (total 2 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   YearsExperience  35 non-null     float64
 1   Salary          35 non-null     int64   
dtypes: float64(1), int64(1)
memory usage: 692.0 bytes
```

In [114]: `df.describe()`

Out[114]:

| | YearsExperience | Salary |
|-------|-----------------|---------------|
| count | 35.000000 | 35.000000 |
| mean | 6.308571 | 83945.600000 |
| std | 3.618610 | 32162.673003 |
| min | 1.100000 | 37731.000000 |
| 25% | 3.450000 | 57019.000000 |
| 50% | 5.300000 | 81363.000000 |
| 75% | 9.250000 | 113223.500000 |
| max | 13.500000 | 139465.000000 |

In [116]: `df.shape`

Out[116]: (35, 2)

In [118]: `df.size`

Out[118]: 70

In [120]: `df.ndim`

Out[120]: 2

In [122]: data.isnull()

Out[122]:

| | YearsExperience | Salary |
|----|-----------------|--------|
| 0 | False | False |
| 1 | False | False |
| 2 | False | False |
| 3 | False | False |
| 4 | False | False |
| 5 | False | False |
| 6 | False | False |
| 7 | False | False |
| 8 | False | False |
| 9 | False | False |
| 10 | False | False |
| 11 | False | False |
| 12 | False | False |
| 13 | False | False |
| 14 | False | False |
| 15 | False | False |
| 16 | False | False |
| 17 | False | False |
| 18 | False | False |
| 19 | False | False |
| 20 | False | False |
| 21 | False | False |
| 22 | False | False |
| 23 | False | False |
| 24 | False | False |
| 25 | False | False |
| 26 | False | False |
| 27 | False | False |
| 28 | False | False |
| 29 | False | False |
| 30 | False | False |
| 31 | False | False |
| 32 | False | False |
| 33 | False | False |
| 34 | False | False |

```
In [124]: df.isnull().any()
```

```
Out[124]: YearsExperience    False
Salary                      False
dtype: bool
```

```
In [126]: df.isnull().sum()
```

```
Out[126]: YearsExperience    0
Salary                      0
dtype: int64
```

```
In [128]: a="ashish"
```

```
In [130]: print(a)
```

```
ashish
```

```
In [132]: a[0]
```

```
Out[132]: 'a'
```

```
In [134]: a[-1]
```

```
Out[134]: 'h'
```

```
In [136]: a[1:3]
```

```
Out[136]: 'sh'
```

```
In [138]: a[1:4]
```

```
Out[138]: 'shi'
```

```
In [150]: #Assigning values in X & Y
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
```

```
#X = df['YearsExperience']
#y = df['Salary']
```

In [154]: `print(X)`

```
[[ 1.1]
 [ 1.3]
 [ 1.5]
 [ 2. ]
 [ 2.2]
 [ 2.9]
 [ 3. ]
 [ 3.2]
 [ 3.2]
 [ 3.7]
 [ 3.9]
 [ 4. ]
 [ 4. ]
 [ 4.1]
 [ 4.5]
 [ 4.9]
 [ 5.1]
 [ 5.3]
 [ 5.9]
 [ 6. ]
 [ 6.8]
 [ 7.1]
 [ 7.9]
 [ 8.2]
 [ 8.7]
 [ 9. ]
 [ 9.5]
 [ 9.6]
 [10.3]
 [10.5]
 [11.2]
 [11.5]
 [12.3]
 [12.9]
 [13.5]]
```

In [156]: `print(y)`

```
[ 39343  46205  37731  43525  39891  56642  60150  54445  64445  57189
  63218  55794  56957  57081  61111  67938  66029  83088  81363  93940
  91738  98273 101302 113812 109431 105582 116969 112635 122391 121872
 127345 126756 128765 135675 139465]
```

In [158]: `import matplotlib.pyplot as plt`
`import seaborn as sns`
`import numpy as np`

In [160]: `#Splitting testdata into X_train,X_test,y_train,y_test`
`from sklearn.model_selection import train_test_split`
`X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=.3,random_st`

```
In [162]: print(X_train)
```

```
[[12.9]
 [ 1.1]
 [ 2.2]
 [ 5.3]
 [ 9.6]
 [ 2.9]
 [ 4. ]
 [ 1.3]
 [ 1.5]
 [12.3]
 [ 2. ]
 [11.2]
 [ 8.2]
 [11.5]
 [ 3.9]
 [ 7.9]
 [ 5.9]
 [ 9. ]
 [ 3. ]
 [ 6.8]
 [13.5]
 [ 3.2]
 [ 4.5]
 [10.3]]
```

```
In [164]: print(X_test)
```

```
[[ 9.5]
 [ 4.1]
 [ 8.7]
 [ 7.1]
 [ 4.9]
 [10.5]
 [ 6. ]
 [ 4. ]
 [ 3.2]
 [ 5.1]
 [ 3.7]]
```

```
In [166]: print(y_train)
```

```
[135675  39343  39891  83088 112635  56642  55794  46205  37731 128765
  43525 127345 113812 126756  63218 101302  81363 105582  60150  91738
 139465  54445  61111 122391]
```

```
In [168]: print(y_test)
```

```
[116969  57081 109431  98273  67938 121872  93940  56957  64445  66029
 57189]
```

```
In [170]: from sklearn.linear_model import LinearRegression  
lr = LinearRegression()  
lr.fit(X_train, y_train)
```

Out[170]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [172]: #Assigning Coefficient (slope) to m  
m = lr.coef_
```

```
In [174]: print("Coefficient :", m)
```

Coefficient : [8555.33918938]

```
In [176]: #Assigning Y-intercept to a  
c = lr.intercept_
```

```
In [178]: #Assigning Y-intercept to a  
c = lr.intercept_
```

Intercept : 29602.07353482097

```
In [180]: lr.score(X_test,y_test) * 100
```

Out[180]: 91.71426108885095

```
In [ ]:
```