

# Logistic Regression

```
In [3]: #Name: Yash Pravin Gadbail  
#Roll no. : 35  
#Sec: 3rd A  
#Sub : ET 1  
#Date:05/10/2024
```

```
In [5]: #Aim: To Perform Operation on Logistic Regression
```

```
In [7]: import pandas as pd  
import matplotlib.pyplot as plt  
import numpy as np  
import seaborn as sns  
from sklearn.model_selection import train_test_split  
import warnings  
warnings.filterwarnings('ignore')
```

```
In [1]: import os
```

```
In [2]: os.getcwd()
```

```
Out[2]: 'C:\\\\Users\\\\RAG'
```

```
In [20]: os.chdir("C:\\\\Users\\\\OneDrive\\\\Desktop")
```

```
In [22]: df=pd.read_csv("C:\\\\Users\\\\OneDrive\\\\Desktop\\\\framingham.csv")
```

```
In [24]: df.head()
```

```
Out[24]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp
0	1	39	4.0	0	0.0	0.0	0	0
1	0	46	2.0	0	0.0	0.0	0	0
2	1	48	1.0	1	20.0	0.0	0	0
3	0	61	3.0	1	30.0	0.0	0	1
4	0	46	3.0	1	23.0	0.0	0	0

In [26]: `df.describe()`

Out[26]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	pre
<b>count</b>	4240.000000	4240.000000	4135.000000	4240.000000	4211.000000	4187.000000	
<b>mean</b>	0.429245	49.580189	1.979444	0.494104	9.005937	0.029615	
<b>std</b>	0.495027	8.572942	1.019791	0.500024	11.922462	0.169544	
<b>min</b>	0.000000	32.000000	1.000000	0.000000	0.000000	0.000000	
<b>25%</b>	0.000000	42.000000	1.000000	0.000000	0.000000	0.000000	
<b>50%</b>	0.000000	49.000000	2.000000	0.000000	0.000000	0.000000	
<b>75%</b>	1.000000	56.000000	3.000000	1.000000	20.000000	0.000000	
<b>max</b>	1.000000	70.000000	4.000000	1.000000	70.000000	1.000000	

In [28]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4240 entries, 0 to 4239
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4240 non-null   int64
1   age                   4240 non-null   int64
2   education             4135 non-null   float64
3   currentSmoker         4240 non-null   int64
4   cigsPerDay            4211 non-null   float64
5   BPMeds                4187 non-null   float64
6   prevalentStroke       4240 non-null   int64
7   prevalentHyp          4240 non-null   int64
8   diabetes              4240 non-null   int64
9   totChol               4190 non-null   float64
10  sysBP                 4240 non-null   float64
11  diaBP                 4240 non-null   float64
12  BMI                   4221 non-null   float64
13  heartRate             4239 non-null   float64
14  glucose               3852 non-null   float64
15  TenYearCHD            4240 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 530.1 KB
```

In [30]: `df.isna().sum()`

```
Out[30]: male                0
age                0
education          105
currentSmoker      0
cigsPerDay         29
BPMeds             53
prevalentStroke    0
prevalentHyp       0
diabetes           0
totChol            50
sysBP              0
diaBP              0
BMI                19
heartRate          1
glucose            388
TenYearCHD         0
dtype: int64
```

In [32]: `df`

```
Out[32]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp
0	1	39	4.0	0	0.0	0.0	0	0
1	0	46	2.0	0	0.0	0.0	0	0
2	1	48	1.0	1	20.0	0.0	0	0
3	0	61	3.0	1	30.0	0.0	0	0
4	0	46	3.0	1	23.0	0.0	0	0
...	...	...	...	...	...	...	...	...
4235	0	48	2.0	1	20.0	NaN	0	0
4236	0	44	1.0	1	15.0	0.0	0	0
4237	0	52	2.0	0	0.0	0.0	0	0
4238	1	40	3.0	0	0.0	0.0	0	0
4239	0	39	3.0	1	30.0	0.0	0	0

4240 rows × 9 columns

## # Missing Value Treatment

In [35]: `df['glucose'].fillna(value = df['glucose'].mean(),inplace=True)`

In [37]: `df['education'].fillna(value = df['education'].mean(),inplace=True)`

In [52]: `df['heartRate'].fillna(value = df['heartRate'].mean(),inplace=True)`

```
In [41]: df['BMI'].fillna(value = df['BMI'].mean(),inplace=True)
```

```
In [45]: df['cigsPerDay'].fillna(value = df['cigsPerDay'].mean(),inplace=True)
```

```
In [48]: df['totChol'].fillna(value = df['totChol'].mean(),inplace=True)
```

```
In [50]: df['BPMeds'].fillna(value = df['BPMeds'].mean(),inplace=True)
```

```
In [54]: df.isna().sum()
```

```
Out[54]: male                0
age                0
education          0
currentSmoker      0
cigsPerDay         0
BPMeds             0
prevalentStroke    0
prevalentHyp       0
diabetes           0
totChol            0
sysBP              0
diaBP              0
BMI                0
heartRate          0
glucose            0
TenYearCHD         0
dtype: int64
```

```
In [56]: #Splitting the dependent and independent variables.
x = df.drop("TenYearCHD",axis=1)
y = df['TenYearCHD']
```

```
In [60]: x
```

```
Out[60]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalent
0	1	39	4.0	0	0.0	0.000000	0	
1	0	46	2.0	0	0.0	0.000000	0	
2	1	48	1.0	1	20.0	0.000000	0	
3	0	61	3.0	1	30.0	0.000000	0	
4	0	46	3.0	1	23.0	0.000000	0	
...	...	...	...	...	...	...	...	
4235	0	48	2.0	1	20.0	0.029615	0	
4236	0	44	1.0	1	15.0	0.000000	0	
4237	0	52	2.0	0	0.0	0.000000	0	
4238	1	40	3.0	0	0.0	0.000000	0	
4239	0	39	3.0	1	30.0	0.000000	0	

4240 rows × 15 columns

# Train Test Split

```
In [63]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,train_si
```

```
In [65]: y_train
```

```
Out[65]: 2399    0
          3609    0
          3814    0
          1941    0
          4021    0
          1878    1
           93    0
          4221    1
          3095    0
          1792    0
           966    0
          2568    0
          3300    0
           818    0
          1541    0
           331    0
          1394    0
           583    0
           866    0
          3148    1
          2890    0
          1884    0
          1279    0
          2581    0
          2326    0
          1999    0
          2133    0
          4032    0
          4168    0
          1801    0
          3874    1
          1709    0
          3126    0
          2664    0
          2616    0
          2253    0
          1147    0
          1794    0
           42    1
          3538    0
          1398    0
          2417    0
          Name: TenYearCHD, dtype: int64
```

## Logistic Regression Algorithm

```
In [68]: from sklearn.linear_model import LogisticRegression  
model = LogisticRegression().fit(x_train,y_train)  
model.score(x_train, y_train)
```

Out[68]: 0.9047619047619048

In [ ]: